

Assignment 1
COMP 409
Marco Guida - 260803123

1.

w = 1000

h = 1000

k = 2000

For n = 1 (single-threaded):

time1: 1971ms

time2: 1863ms

time3: 1910ms

time4: 1806ms

time5: 2073ms

average: 1924.6ms

For n = 2:

time1: 1466ms

time2: 1556ms

time3: 1482ms

time4: 1559ms

time5: 1478ms

average: 1508.2ms

speedup: 1.2760907

For n = 3:

time1: 1010ms

time2: 1223ms

time3: 1024ms

time4: 1006ms

time5: 1233ms

average: 1099.2ms

speedup: 1.75090975

For n = 4:

time1: 1042ms

time2: 789ms


time3: 1010ms

time4: 1016ms

time5: 825ms

average: 936.4ms

speedup: 2.05531824



I observed a speedup with each additional thread. As expected, the magnitude of the speedup increase gets smaller with each additional thread (apart from the speedup from two threads to three threads due to the noticeable outlier that occurred when the program used two threads) as speedup is limited by that part of the program that cannot be parallelized. The parts of the program that cannot be parallelized, and thus cause the speedup curve above, are the blocks of instructions that read/write the shared variable `k`, which represents the number of rectangles to draw, and `pixelsUsed`, which monitors whether a pixel is currently being modified by another thread. Reading/writing to `pixelsUsed` cannot be parallelized to ensure that a thread that is drawing a new rectangle does not overlap other in-progress rectangles, and so it may take longer to find a suitable location for a thread to draw a rectangle.